



Written examination, December 7, 2016 Page 1 of 10 pages

Course name: Mathematical software programming

Course number: 02635

Aids allowed: All aids allowed

Exam duration: 4 hours

Weighting: 80/100

Final exam
Mathematical Software Programming

This exam contains a total of 20 questions: 16 multiple choice questions (questions 1–16) and 4 programming questions (questions 17–20). Your exam answers must be submitted electronically as a **PDF document**. You may include your code in the document along with your answers or submit the code separately in a ZIP file.

1. (2 points) The associative property

$$(a + b) + c = a + (b + c)$$

always holds for

- (a) floating point arithmetic with floating point numbers a , b , and c (e.g. `double`)

- A. True
- B. False

- (b) integer arithmetic with integers a , b , and c (e.g. `int`)

- A. True
- B. False

2. (2 points) Consider the following system of equations with two unknowns x and y

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} e \\ f \end{bmatrix}$$

where a , b , c , d , e , and f are real numbers. Recall that if the system matrix is nonsingular (i.e., the determinant is nonzero), the solution can be expressed as

$$\begin{bmatrix} x \\ y \end{bmatrix} = \frac{1}{ad - cb} \begin{bmatrix} d & -b \\ -c & a \end{bmatrix} \begin{bmatrix} e \\ f \end{bmatrix}.$$

Catastrophic cancellation is likely to occur in the above computations if

- A. $|ad - cb|$ is large
- B. $|de - bf|$ or $|af - ce|$ is large
- C. $a = d = 1$ and $c = b = 10^{-8}$
- D. $ad \approx cb$ and/or $de \approx bf$ and/or $ce \approx af$

3. (4 points) Consider the following nested loops

```
for(i=0; i<M; i++)      // outer loop
  for(j=0; j<N; j++)    // inner loop
    a[i] = j*a[i];
```

where **a** is an array of length M.

- (a) References to **a** are temporally local with respect to the outer loop.
- A. True
 - B. False
- (b) References to **a** are spatially local with respect to the outer loop.
- A. True
 - B. False
4. (2 points) Cache memory is generally slower than main memory (RAM).
- A. True
 - B. False
5. (2 points) When parallelizing a program, the speedup of the execution time using p cores is defined as
- $$S(p) = \frac{T(1)}{T(p)}$$
- where $T(p)$ denotes the execution time with p cores. What is the maximum speedup that is possible using any number of cores if only 80% of the program can be parallelized?
- A. 1.25
 - B. 1.6
 - C. 4
 - D. 5
6. (2 points) Memory leaks are possible in C but not in C++.
- A. True
 - B. False

7. (6 points) Consider the recursive function definition

$$f(0) = 0, \quad f(n) = f(n-1) + 2n - 1, \quad n \geq 1,$$

where n is a positive integer. This function can be implemented in C as follows:

```
unsigned long fun(unsigned long n) {  
    if (n == 0) return 0;  
    return fun(n-1) + 2*n - 1;  
}
```

- (a) What type of recursion is this?
- A. Single recursion
 - B. Double recursion
 - C. Multiple recursion
- (b) What is the time complexity of this implementation?
- A. $O(1)$
 - B. $O(n)$
 - C. $O(n^2)$
 - D. $O(2^n)$
- (c) What is the space complexity of this implementation?
- A. $O(1)$
 - B. $O(n)$
 - C. $O(n^2)$
 - D. $O(2^n)$
8. (2 points) Templates in C++ facilitate
- A. functional programming
 - B. object-oriented programming
 - C. generic programming
 - D. specific programming

9. (2 points) Consider the following C function that evaluates the function

$$f(x) = ax^2 + bx + c$$

```
double * quad_eval(double a, double b, int double c, double x) {  
    double *pv;  
    double val = b;  
    pv = &val;  
    val += a*x;  
    val *= x;  
    val += c;  
    return pv;  
}
```

There is a problem with this implementation. What is the problem?

- A. The result is subject to catastrophic cancellation when $|a|$ is large
 - B. The function returns a pointer to a local variable.
 - C. The function returns b instead of $f(x)$
 - D. The function leaks memory
10. (2 points) Dynamic memory allocation is faster than automatic memory allocation.
- A. True
 - B. False
11. (2 points) The operator `::` in C++ (e.g. `std::cout`) is referred to as the
- A. member access operator
 - B. scope resolution operator
 - C. namespace operator
 - D. function operator

12. (2 points) Consider the following C program:

```
#include <stdlib.h>
#include <stdio.h>
#define N 50

double * init_array(size_t n, double h) {
    double * p = malloc(n*sizeof(*p));
    if (p != NULL) {
        for (size_t i=0; i<n; i++)
            *(p++) = i*h;
    }
    return p;
}

int main(void) {

    double *p = init_array(N);
    if (p == NULL) return EXIT_FAILURE;

    for (size_t k=0; k<N; k++) printf("%3d %6g\n", k, p[k]);

    free(p);
    return 0;
}
```

Is there a problem with this program? If yes, what is it?

- A. Yes, `init_array` returns a local variable.
- B. Yes, the return value of `init_array` does not point to the first element of the array.
- C. Yes, `init_array` should call `free` before returning to prevent a memory leak.
- D. No, the program is fine.

13. (2 points) When could it be necessary to know the endianness of a computer?

- A. When sharing binary data or files between different computers.
- B. When reading input from the user (e.g. with `scanf`).
- C. When reading or writing numbers from/to a text file.
- D. When allocating memory for multi-byte data types.

14. (6 points) Recall that a floating point number x can be represented as

$$x = s \cdot (d_0.d_1d_2\dots d_n)_b \cdot b^e = s \cdot \sum_{i=0}^n d_i b^{e-i}$$

where s denotes the sign, e is the exponent, $b \geq 2$ is the base, d_0, \dots, d_n comprise the mantissa and $d_i \in \{0, 1, \dots, b-1\}$.

- (a) What characterizes a *normal* floating point number?

- A. $b = 1$
- B. $e = 0$
- C. $s = +1$
- D. $d_0 \neq 0$

- (b) What is the smallest representable number greater than 1.0?

- A. $1 + b^{-(n-1)}$
- B. $1 + b^{-n}$
- C. $1 + b^{-(n+1)}$
- D. $1 + b^{-(n+1)}/2$

- (c) What is the largest representable number less than 1.0?

- A. $1 - b^{-(n-1)}$
- B. $1 - b^{-n}$
- C. $1 - b^{-(n+1)}$
- D. $1 - b^{-(n+1)}/2$

15. (2 points) Consider the following piece of code:

```
double data[4] = {5.0, -2.0, 2.0, 0.0};  
data[1] *= 2.0;  
--data[1];
```

The first element of the array is initialized as 5.0. What is the value of this element of the array `data` after running the code?

- A. -5.0
- B. 5.0
- C. 9.0
- D. 10.0

16. (2 points) Consider the following C struct and array:

```
typedef struct matrix {  
    size_t m;  
    size_t n;  
    double **A;  
} matrix_t;  
  
matrix_t mat_arr[10];
```

Which of the following options is a correct way of assigning the value 4 to the field `m` of the second element of the array `mat_arr`?

- A. `(mat_arr+1).m = 4;`
 - B. `*mat_arr[1].m = 4;`
 - C. `mat_arr[1]->m = 4;`
 - D. `(mat_arr+1)->m = 4;`
 - E. `*(mat_arr+1)->m = 4;`
17. (12 points) The cumulative sum of the elements of a vector $x = (x_1, \dots, x_n)$ can be expressed as

$$y_i = \sum_{k=1}^i x_k, \quad i = 1, \dots, n.$$

- (a) Implement a function that computes the cumulative sum. Your function should have the following prototype:

```
void cumsum(  
    const double * x,    /* pointer to x[0] */  
    double * y,          /* pointer to y[0] */  
    unsigned long n      /* length of arrays */  
);
```

The inputs `x` and `y` should be pointers to the first element of two non-overlapping arrays (corresponding to x and y), and the third input `n` represents the length of the arrays.

- (b) What is the time complexity of computing the cumulative sum?
- (c) Explain how you tested your function and/or include your test program(s).

18. (4 points) Recall that a sparse matrix can be stored in triple form which can be represented using the following abstract data type:

```
typedef struct sparse_triplet {
    unsigned long m; /* number of rows */
    unsigned long n; /* number of columns */
    unsigned long nnz; /* number of nonzeros */
    unsigned long * I; /* pointer to array with row indices */
    unsigned long * J; /* pointer to array with column indices */
    double * V; /* pointer to array with values */
} sparse_triplet_t;
```

Write a function that transposes such a sparse matrix in-place (i.e., without allocating a new `sparse_triplet_t`). Your function should have the following prototype:

```
void trans_sparse_triplet(sparse_triplet_t * sp);
```

19. (12 points) An exponential moving average of a series of numbers x_0, x_1, x_2, \dots is defined recursively as

$$y_0 = x_0, \quad y_k = \alpha x_k + (1 - \alpha)y_{k-1}, \quad k \geq 1$$

where $\alpha \in (0, 1)$ is a constant.

- (a) Write a function that computes an exponential moving average of a series of numbers, x_0, x_1, \dots, x_{n-1} , stored as an array of length n . Use the following prototype:

```
void ema(
    const double * x, /* pointer to x[0] */
    double * y, /* pointer to y[0] */
    unsigned long n, /* length of arrays */
    double alpha /* constant alpha */
);
```

- (b) What is the time complexity of computing an exponential moving average?
- (c) Explain how you tested your function and/or include your test program(s).

20. (10 points) Two vectors x and y of length n are orthogonal if and only if their inner product

$$x^T y = \sum_{i=1}^n x_i y_i$$

is equal to zero.

- (a) Write a function that checks if two vector are orthogonal. Use the following prototype:

```
int orthogonal(  
    const double * x,      /* pointer to x[0]    */  
    const double * y,      /* pointer to y[0]    */  
    unsigned long n        /* length of arrays   */  
);
```

The inputs x and y should be pointers to the first element of two arrays of length n . The output should be equal to 1 if the vectors are orthogonal, and otherwise the output should be 0.

- (b) Test your function with the following vectors as input:

$$x = \begin{pmatrix} 0.1 \\ 0.3 \end{pmatrix}, \quad y = \begin{pmatrix} 3 \\ -1 \end{pmatrix}.$$

Does your function work as expected? Why/why not?