

Technical University of Denmark

Written examination, December 10, 2019

Page 1 of 10 pages

Course title:	Mathematical Software Programming
---------------	-----------------------------------

Course number:	02635
----------------	-------

Aids allowed:	All aids allowed
---------------	------------------

Exam duration:	4 hours
----------------	---------

Weighting:	80/100
------------	--------

---

**Final exam**  
**Mathematical Software Programming**

This exam contains a total of 16 questions: 12 multiple choice questions (questions 1–12) and 4 programming questions (questions 13–16). Your exam answers to the multiple choice questions must be submitted electronically as a **PDF document**, and your code must be attached (e.g., as a **ZIP file**).

---

1. (4 points) The unit round-off for double precision floating point numbers is  $u = 2^{-53}$ .
  - (a) The unit round-off  $u$  can be represented as a double precision floating point number.
    - A. True
    - B. False
  - (b) The number  $1 + u$  can be represented as a double precision floating point number.
    - A. True
    - B. False
2. (4 points) Consider the real-valued function  $f(x) = ax^2 + bx + c$  where  $a, b, c, x \in \mathbb{R}$ .
  - (a) What is the relative condition number of  $f$ ?
    - A.  $c_f(x) = \frac{|ax^2+bx+c|}{|2ax+b|}$
    - B.  $c_f(x) = \frac{|2ax+b|}{|ax^2+bx+c|}$
    - C.  $c_f(x) = \frac{|2ax^2+bx|}{|ax^2+bx+c|}$
    - D.  $c_f(x) = \frac{|ax^2+bx+c|}{|2ax^2+bx|}$
  - (b) Suppose  $a = 1$ ,  $b = 4$ , and  $c = 4$ . For what values of  $x$  is the problem of evaluating  $f(x)$  ill-posed?
    - A. When  $x = 0$ .
    - B. When  $x = -1$ .
    - C. When  $x = -2$ .
    - D. When  $x \in \{-2, 0\}$ .

3. (6 points) Let  $f(x) = \log(1+x)$  where  $x$  is a positive real number.

(a) The relative condition number of  $f$  is given by

A.  $c_f(x) = \left| \frac{1}{x \log(1+x)} \right|$

B.  $c_f(x) = \left| \frac{1}{(1+x) \log(1+x)} \right|$

C.  $c_f(x) = \left| \frac{1+x}{x \log(1+x)} \right|$

D.  $c_f(x) = \left| \frac{x}{(1+x) \log(1+x)} \right|$

(b) What can be said about the relative condition number of  $f$  when  $|x|$  is close to 0?

A.  $c_f(x)$  is close to 0

B.  $c_f(x)$  is close to 1

C.  $c_f(x)$  is much larger than 1

(c) Consider the following implementation of the function  $f(x) = \log(1+x)$ :

```
#include <math.h>
double fun(double x) {
    return log(1.0+x);
}
```

This implementation may result in a large relative error when  $|x|$  is close to 0. Why?

A. The condition number of  $\log(x)$  is large when  $x$  is close to 1.

B. The sum  $1+x$  is prone to catastrophic cancellation when  $|x|$  is close to 0.

C. The sum  $1+x$  may underflow when  $|x|$  is close to 0.

4. (4 points) The Chebyshev polynomials of the first kind can be defined recursively as

$$T_0(x) = 1$$

$$T_1(x) = x$$

$$T_k(x) = 2xT_{k-1}(x) - T_{k-2}(x), \quad k = 2, 3, \dots$$

- (a) What kind of recursion is this?

- A. single recursion
- B. multiple recursion

- (b) What is the complexity associated with the following implementation?

```
double T(double x, unsigned int k) {  
    if (k==0) return 1.0;  
    if (k==1) return x;  
    return 2*x*T(x,k-1)-T(x,k-2);  
}
```

- A.  $O(k)$  space and  $O(k)$  time
- B.  $O(2^k)$  space and  $O(k)$  time
- C.  $O(k)$  space and  $O(2^k)$  time
- D.  $O(2^k)$  space and  $O(2^k)$  time

5. (4 points) Suppose the variable **A** represents a two-dimensional array of size  $m \times n$ . Consider the following code segment:

```
for (int i=0; i<m; i++) {  
    for (int j=0; j<n; j++) {  
        y[i] += A[i][j]*x[j];  
    }  
}
```

- (a) References to **y[i]** are temporally local.

- A. True
- B. False

- (b) References to the array represented by **x** are spatially local.

- A. True
- B. False

6. (6 points) The theoretical improvement in speed of execution of a task executed on  $p$  processors can be expressed as

$$S(p) = \frac{T(1)}{T(p)} = \frac{fT(1) + (1-f)T(1)}{(f/p)T(1) + (1-f)T(1)}$$

where  $T(p)$  is the execution time on  $p$  processors (real time) and  $f$  is the so-called parallel fraction of the task. For example, if 50% of a task can be parallelized, then  $f = 0.5$ .

Suppose that a specific task consists of two subtasks. The first subtask retrieves data from a file; it takes 20 seconds to execute on a given system, and it cannot be parallelized. The second subtask processes a large number of data records; it takes 60 seconds to execute sequentially. Each data record can be processed independently, so the second subtask is easily parallelized.

- (a) What is the parallel fraction of the task consisting of the two subtasks?

- A.  $f = 1/4$
- B.  $f = 1/3$
- C.  $f = 3/4$
- D.  $f = 4/5$

- (b) What is the theoretical speedup for  $p = 9$  processors?

- A. 3
- B. 4
- C. 6
- D. 9


7. (2 points) Memory leaks are possible in C but not in C++.

- A. True
- B. False

8. (2 points) The operator `::` in C++ (e.g. `std::cout`) is referred to as the

- A. member access operator
- B. scope resolution operator
- C. namespace operator
- D. function operator

9. (4 points) The following plot shows the CPU time required by some algorithm to solve a certain problem as a function of its dimension  $n$ .



What is the time complexity of the algorithm?

- A.  $O(n)$
  - B.  $O(n^{3/2})$
  - C.  $O(n^2)$
  - D.  $O(2^n)$
10. (4 points) Consider the following function:

```
double * ones(size_t n) {  
    double *arr = malloc(n*sizeof(*arr));  
    if (arr != NULL) {  
        for (size_t k=0; k<n; k++)  
            arr[k] = 1.0;  
    }  
    return arr;  
}
```

- (a) What kind of memory allocation is used to allocate what `arr` points to?
- A. Automatic memory allocation
  - B. Dynamic memory allocation
- (b) What kind of memory allocation is used to allocate the pointer `arr`?
- A. Automatic memory allocation
  - B. Dynamic memory allocation

11. (4 points) Suppose the matrix

$$A = \begin{bmatrix} A_{11} & A_{12} & A_{13} & A_{14} \\ A_{21} & A_{22} & A_{23} & A_{24} \\ A_{31} & A_{32} & A_{33} & A_{34} \end{bmatrix}$$

is stored as a *one-dimensional, row-major* **double** array of length 12.

The elements of an array of double-precision floating-point numbers can be scaled using the CBLAS routine `cblas_dscal()` which has the following prototype:

```
void cblas_dscal(
    const int n,           // number of elements
    const double a,        // scalar a
    double * x,            // pointer to first element of array
    const int incx         // stride
);
```

Now suppose `pA` is a pointer to  $A_{11}$  (i.e., `pA` is a **double \***).

(a) Which of the following calls to `cblas_dscal()` scales the elements  $A_{22}$ ,  $A_{23}$ , and  $A_{24}$  by  $-1$ ?

- A. `cblas_dscal(3, -1.0, pA+4, 1)`
- B. `cblas_dscal(3, -1.0, pA+4, 3)`
- C. `cblas_dscal(3, -1.0, pA+4, 4)`
- D. `cblas_dscal(3, -1.0, pA+5, 1)`
- E. `cblas_dscal(3, -1.0, pA+5, 3)`
- F. `cblas_dscal(3, -1.0, pA+5, 4)`

(b) Consider the following `cblas_dscal()` call:

```
cblas_dscal(4, 2.0, pA+1, 2);
```

Which elements of  $A$  are being scaled by a factor of 2?

- A.  $A_{11}, A_{13}, A_{21}, A_{23}$
- B.  $A_{12}, A_{14}, A_{22}, A_{24}$
- C.  $A_{21}, A_{12}, A_{32}, A_{23}$
- D.  $A_{11}, A_{31}, A_{22}, A_{13}$

12. (2 points) Cache memory is slower than main memory.

- A. True
- B. False

13. (8 points) Implement a function that performs the matrix-vector multiplication

$$y \leftarrow A^T x + y$$

where  $A$  is a sparse matrix of size  $m \times n$  in triplet form, and  $x \in \mathbb{R}^m$  and  $y \in \mathbb{R}^n$  are vectors.

### Requirements

- Your function must have the following prototype:

```
void sparse_triplet_mv_trans(  
    const struct sparse_triplet *A,  
    const double *x,  
    double *y);
```

You may assume that all inputs are valid.

- The sparse matrix  $A$  should be stored using the sparse triplet format represented by the following data structure:

```
/* Structure representing a sparse matrix in triplet form */  
struct sparse_triplet {  
    size_t m;    /* number of rows */  
    size_t n;    /* number of columns */  
    size_t nnz;  /* number of nonzeros */  
    size_t * I;  /* pointer to array with row indices */  
    size_t * J;  /* pointer to array with column indices */  
    double * V;  /* pointer to array with values */  
};
```

The header file `sparse_triplet.h` in the ZIP file attached to the exam defines the `sparse_triplet` data structure. You do not need to include this header file when you submit your solution.

- Use the template `sparse_triplet_mv_trans.c` for your implementation. The template is included in the ZIP file attached to the exam.



14. (8 points) The Poisson distribution is a discrete probability distribution with probability mass function

$$f(k; \lambda) = \frac{\lambda^k}{k!} e^{-\lambda}, \quad k \in \mathbb{N}_0,$$

where the parameter  $\lambda > 0$  is the so-called rate. Implement a function that evaluates  $f(k; \lambda)$ .

### Requirements

- Your function must have the following prototype:

```
double poisson_pmf(unsigned long k, double lambda);
```

- The function should return NAN if  $\lambda$  is not positive.
- Use the template `poisson_pmf.c` for your implementation. The template is included in the ZIP file attached to the exam.

15. (8 points) The centering matrix of order  $n$  is given by

$$C = I - \frac{1}{n} \mathbf{1} \mathbf{1}^T$$

where  $I$  is the  $n \times n$  identity matrix and  $\mathbf{1}$  is the vector of length  $n$  whose entries are all equal to 1. For example, the centering matrix of order 3 is the matrix

$$C = \begin{bmatrix} \frac{2}{3} & -\frac{1}{3} & -\frac{1}{3} \\ -\frac{1}{3} & \frac{2}{3} & -\frac{1}{3} \\ -\frac{1}{3} & -\frac{1}{3} & \frac{2}{3} \end{bmatrix}.$$

Implement a function that computes  $x \leftarrow Cx$ , i.e., the function should overwrite a vector  $x \in \mathbb{R}^n$  by the matrix-vector product  $Cx$ .

### Requirements

- The function must have the following prototype:

```
int dcmv(int n, double * x);
```

The input `n` is the order of the matrix  $C$ , and `x` is a pointer to the first elements of an array of length `n` that represent the vector  $x$ .

- The function should return the value -1 in case of an error or invalid input, and otherwise it should return the value 0.
- Use the template `dcmv.c` for your implementation. The template is included in the ZIP file attached to the exam.

16. (10 points) Consider the system of equations

$$\begin{bmatrix} c & -s \\ s & c \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} = \begin{bmatrix} r \\ 0 \end{bmatrix} \quad (1)$$

where  $r \geq 0$ ,  $c = \cos(\theta)$ , and  $s = \sin(\theta)$  for some  $\theta \in [0, 2\pi]$ . It is easy to check that  $a, b$  must satisfy

$$\begin{bmatrix} a \\ b \end{bmatrix} = r \begin{bmatrix} c \\ -s \end{bmatrix},$$

and hence  $r = \sqrt{a^2 + b^2}$ . It follows that if  $r > 0$ , then  $c = a/r$  and  $s = -b/r$  (the angle  $\theta$  can be chosen arbitrarily when  $r = 0$ , e.g.,  $c = 1$  and  $s = 0$ ).

- (a) Implement a function that takes  $a$  and  $b$  as inputs and computes  $c$  and  $s$  such that  $a$  and  $b$  satisfies (1).

### Requirements

- Your function must have the following prototype:

```
void rotg(double a, double b, double * c, double * s);
```

The inputs  $c$  and  $s$  point to the locations where  $c$  and  $s$  should be stored. You may assume that all inputs are valid.

- Use the template `rotg.c` for your implementation. The template is included in the ZIP file attached to the exam.

- (b) Implement a function that applies the transformation

$$\begin{bmatrix} x_i \\ y_i \end{bmatrix} \leftarrow \begin{bmatrix} c & -s \\ s & c \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix}, \quad i = 1, \dots, n,$$

where  $x = (x_1, \dots, x_n)$  and  $y = (y_1, \dots, y_n)$  are vectors of length  $n$ , and  $c = \cos(\theta)$  and  $s = \sin(\theta)$  for some  $\theta \in [0, 2\pi]$ .

### Requirements

- Your function must have the following prototype:

```
int rot(double c, double s, int n, double * x, double * y);
```

The input  $n$  is the length of the vectors  $x$  and  $y$ . The inputs  $x$  and  $y$  are pointers to the first element of the arrays that represent  $x$  and  $y$ , respectively.

- The function should return the value  $-1$  in case of an invalid input, and otherwise it should return the value  $0$ .
- Use the template `rot.c` for your implementation. The template is included in the ZIP file attached to the exam.